

```
from pyspark.sql import SparkSession
from pyspark.sql.functions import col
spark = SparkSession.builder \
    .appName('online shopping platform') \
    .getOrCreate()
```

```
from pyspark.sql.types import *
from pyspark.sql.functions import *
```

## ▼ PHASE 1 — DATA INGESTION & SCHEMA

```
orders_data = [
    ("0001", "Delhi ", "Laptop", "45000", "2024-01-05", "Completed"),
    ("0002", "Mumbai", "Mobile ", "32000", "05/01/2024", "Completed"),
    ("0003", "Bangalore", "Tablet", "30000", "2024/01/06", "Completed"),
    ("0004", "Delhi", "Laptop", "", "2024-01-07", "Cancelled"),
    ("0005", "Mumbai", "Mobile", "invalid", "2024-01-08", "Completed"),
    ("0006", "Chennai", "Tablet", None, "2024-01-08", "Completed"),
    ("0007", "Delhi", "Laptop", "47000", "09-01-2024", "Completed"),
    ("0008", "Bangalore", "Mobile", "28000", "2024-01-09", "Completed"),
    ("0009", "Mumbai", "Laptop", "55000", "2024-01-10", "Completed"),
    ("0009", "Mumbai", "Laptop", "55000", "2024-01-10", "Completed")
]
columns = ["order_id", "city ", "product", "price", "order_date", "status"]
```

### 1. Define an explicit schema

```
orders_schema = StructType([
    StructField("order_id", StringType(), True),
    StructField("city", StringType(), True),
    StructField("product", StringType(), True),
    StructField("price", StringType(), True),
    StructField("order_date", StringType(), True),
    StructField("status", StringType(), True)
])
```

### 2. Create a DataFrame using the schema

```
orders_df = spark.createDataFrame(data=orders_data, schema=orders_schema)
```

Print schema and validate data types

```
orders_df.printSchema()
orders_df.show(truncate=False)
```

```
root
|-- order_id: string (nullable = true)
|-- city: string (nullable = true)
|-- product: string (nullable = true)
|-- price: string (nullable = true)
|-- order_date: string (nullable = true)
|-- status: string (nullable = true)

+-----+-----+-----+-----+-----+
|order_id|city      |product  |price    |order_date|status   |
+-----+-----+-----+-----+-----+
|0001    |Delhi     |Laptop   |45000   |2024-01-05|Completed|
|0002    |Mumbai    |Mobile   |32000   |05/01/2024|Completed|
|0003    |Bangalore |Tablet   |30000   |2024/01/06|Completed|
|0004    |Delhi     |Laptop   |        |2024-01-07|Cancelled|
|0005    |Mumbai    |Mobile   |invalid |2024-01-08|Completed|
|0006    |Chennai   |Tablet   |NULL    |2024-01-08|Completed|
|0007    |Delhi     |Laptop   |47000   |09-01-2024|Completed|
|0008    |Bangalore |Mobile   |28000   |2024-01-09|Completed|
|0009    |Mumbai    |Laptop   |55000   |2024-01-10|Completed|
|0009    |Mumbai    |Laptop   |55000   |2024-01-10|Completed|
+-----+-----+-----+-----+-----+
```

## ▼ PHASE 2 – DATA CLEANING

### 4. Trim all string columns

```
orders_df = orders_df.withColumn("city", lower(trim(col("city")))) \
                     .withColumn("product", lower(trim(col("product"))))
```

```
orders_df.show(truncate=False)
orders_df.printSchema()
```

```
+-----+-----+-----+-----+-----+
|order_id|city      |product  |price    |order_date|status   |
+-----+-----+-----+-----+-----+
|0001    |delhi     |laptop   |45000   |2024-01-05|Completed|
|0002    |mumbai    |mobile   |32000   |05/01/2024|Completed|
|0003    |bangalore |tablet   |30000   |2024/01/06|Completed|
|0004    |delhi     |laptop   |        |2024-01-07|Cancelled|
|0005    |mumbai    |mobile   |invalid |2024-01-08|Completed|
|0006    |chennai   |tablet   |NULL    |2024-01-08|Completed|
|0007    |delhi     |laptop   |47000   |09-01-2024|Completed|
|0008    |bangalore |mobile   |28000   |2024-01-09|Completed|
|0009    |mumbai    |laptop   |55000   |2024-01-10|Completed|
|0009    |mumbai    |laptop   |55000   |2024-01-10|Completed|
+-----+-----+-----+-----+-----+
```

```
root
|-- order_id: string (nullable = true)
|-- city: string (nullable = true)
|-- product: string (nullable = true)
|-- price: string (nullable = true)
|-- order_date: string (nullable = true)
```

```
|-- status: string (nullable = true)
```

## 5. Standardize city and product values

```
print("Distinct City values after standardization:")
orders_df.select("city").distinct().show(truncate=False)

print("Distinct Product values after standardization:")
orders_df.select("product").distinct().show(truncate=False)

Distinct City values after standardization:
+-----+
|city    |
+-----+
|delhi   |
|bangalore|
|mumbai  |
|chennai |
+-----+

Distinct Product values after standardization:
+-----+
|product|
+-----+
|mobile  |
|tablet  |
|laptop  |
+-----+
```

## 6. Convert amount to IntegerType

```
orders_df = orders_df.withColumn("price", col("price").cast(IntegerType()))
orders_df.printSchema()

root
 |-- order_id: string (nullable = true)
 |-- city: string (nullable = true)
 |-- product: string (nullable = true)
 |-- price: integer (nullable = true)
 |-- order_date: string (nullable = true)
 |-- status: string (nullable = true)
```

## 7. Handle invalid and null amount values

```
orders_df = orders_df.withColumn(
    "price",
    when(col("price").isNull() | (col("price") == "") | (col("price") == "indefinite"))
    .otherwise(col("price"))
    .cast(IntegerType())
)
```

```
orders_df.printSchema()
orders_df.show(truncate=False)
```

```
root
|-- order_id: string (nullable = true)
|-- city: string (nullable = true)
|-- product: string (nullable = true)
|-- price: integer (nullable = true)
|-- order_date: string (nullable = true)
|-- status: string (nullable = true)
```

order_id	city	product	price	order_date	status
0001	delhi	laptop	45000	2024-01-05	Completed
0002	mumbai	mobile	32000	05/01/2024	Completed
0003	bangalore	tablet	30000	2024/01/06	Completed
0004	delhi	laptop	0	2024-01-07	Cancelled
0005	mumbai	mobile	0	2024-01-08	Completed
0006	chennai	tablet	0	2024-01-08	Completed
0007	delhi	laptop	47000	09-01-2024	Completed
0008	bangalore	mobile	28000	2024-01-09	Completed
0009	mumbai	laptop	55000	2024-01-10	Completed
0009	mumbai	laptop	55000	2024-01-10	Completed

## 8. Remove duplicate orders

```
orders_df = orders_df.dropDuplicates()
orders_df.printSchema()
orders_df.show(truncate=False)
```

```
root
|-- order_id: string (nullable = true)
|-- city: string (nullable = true)
|-- product: string (nullable = true)
|-- price: integer (nullable = true)
|-- order_date: string (nullable = true)
|-- status: string (nullable = true)
```

order_id	city	product	price	order_date	status
0002	mumbai	mobile	32000	05/01/2024	Completed
0005	mumbai	mobile	0	2024-01-08	Completed
0001	delhi	laptop	45000	2024-01-05	Completed
0003	bangalore	tablet	30000	2024/01/06	Completed
0008	bangalore	mobile	28000	2024-01-09	Completed
0007	delhi	laptop	47000	09-01-2024	Completed
0009	mumbai	laptop	55000	2024-01-10	Completed
0006	chennai	tablet	0	2024-01-08	Completed

## 9. Keep only Completed orders

```
orders_df = orders_df.filter(col("status") == "Completed")
orders_df.printSchema()
orders_df.show(truncate=False)
```

```
root
|-- order_id: string (nullable = true)
|-- city: string (nullable = true)
|-- product: string (nullable = true)
|-- price: integer (nullable = true)
|-- order_date: string (nullable = true)
|-- status: string (nullable = true)
```

order_id	city	product	price	order_date	status
0002	mumbai	mobile	32000	05/01/2024	Completed
0005	mumbai	mobile	0	2024-01-08	Completed
0001	delhi	laptop	45000	2024-01-05	Completed
0003	bangalore	tablet	30000	2024/01/06	Completed
0008	bangalore	mobile	28000	2024-01-09	Completed
0007	delhi	laptop	47000	09-01-2024	Completed
0009	mumbai	laptop	55000	2024-01-10	Completed
0006	chennai	tablet	0	2024-01-08	Completed

## ▼ PHASE 3 — BASIC ANALYTICS

### 10. Total revenue per city

```
city_revenue_df = orders_df.groupBy("city").agg(sum("price").alias("total_revenue"))
city_revenue_df.show()
```

city	total_revenue
chennai	0
delhi	92000
bangalore	58000
mumbai	87000

### 11. Total revenue per product

```
product_revenue_df = orders_df.groupBy("product").agg(sum("price").alias("total_revenue"))
product_revenue_df.show()
```

product	total_revenue
mobile	60000
tablet	30000
laptop	147000

## 12. Average order value per city

```
avg_order_value_df = orders_df.groupBy("city").agg(avg("price").alias("avg_c"))
avg_order_value_df.show()
```

city	avg_order_value
chennai	0.0
delhi	46000.0
bangalore	29000.0
mumbai	29000.0

## ▼ PHASE 4 — WINDOW FUNCTION

### 13. Rank cities by total revenue

```
from pyspark.sql.window import Window
from pyspark.sql.functions import row_number
window_spec = Window.orderBy(col("total_revenue").desc())
ranked_cities_df = city_revenue_df.withColumn("rank", row_number().over(window_spec))
ranked_cities_df.show()
```

city	total_revenue	rank
delhi	92000	1
mumbai	87000	2
bangalore	58000	3
chennai	0	4

### 14. Identify top-performing city

```
top_city_df = ranked_cities_df.orderBy(col("total_revenue").desc()).limit(1)
top_city_df.show()
```

city	total_revenue	rank
delhi	92000	1

```
+-----+-----+-----+
|delhi|      92000|    1|
+-----+-----+-----+
```

## ▼ PHASE 5 — PERFORMANCE AWARENESS

### 15. Cache the cleaned DataFrame

```
orders_df.cache()
orders_df.count()
```

8

### 16. Run two aggregations and observe behavior

```
orders_df.groupBy("city").agg(sum("price").alias("total_revenue")).show()
```

```
+-----+-----+
|     city|total_revenue|
+-----+-----+
|bangalore|      58000|
|mumbai|      87000|
|delhi|      92000|
|chennai|          0|
+-----+-----+
```

```
orders_df.groupBy("product").agg(sum("price").alias("total_revenue")).show()
```

```
+-----+-----+
|product|total_revenue|
+-----+-----+
| mobile|      60000|
| laptop|     147000|
| tablet|      30000|
+-----+-----+
```

### 17. Use explain(True) to inspect the plan

```
orders_df.explain(True)
```

```
-- Parsed Logical Plan --
'Project [unresolvedstarwithcolumns(price, cast('price as int), None)]
+- Filter (status#5 = Completed)
  +- Deduplicate [city#26, order_id#0, price#47, product#27, order_date#4, s
    +- Filter (status#5 = Completed)
      +- Deduplicate [city#26, order_id#0, price#47, product#27, order_dat
        +- Project [order_id#0, city#26, product#27, cast(CASE WHEN ((isn
          +- Project [order_id#0, city#26, lower(trim(product#2, None))]
```

```

+- Project [order_id#0, lower(trim(city#1, None)) AS city#2]
  +- LogicalRDD [order_id#0, city#1, product#2, price#3, o

== Analyzed Logical Plan ==
order_id: string, city: string, product: string, price: int, order_date: stri
Project [order_id#0, city#26, product#27, cast(price#47 as int) AS price#132,
+- Filter (status#5 = Completed)
  +- Deduplicate [city#26, order_id#0, price#47, product#27, order_date#4, s
    +- Filter (status#5 = Completed)
      +- Deduplicate [city#26, order_id#0, price#47, product#27, order_dat
        +- Project [order_id#0, city#26, product#27, cast(CASE WHEN ((isn
          +- Project [order_id#0, city#26, lower(trim(product#2, None))
            +- Project [order_id#0, lower(trim(city#1, None)) AS city#2
              +- LogicalRDD [order_id#0, city#1, product#2, price#3, o

== Optimized Logical Plan ==
Aggregate [city#26, order_id#0, price#47, product#27, order_date#4, status#5]
+- Project [order_id#0, lower(trim(city#1, None)) AS city#26, lower(trim(prod
  +- Filter (isnotnull(status#5) AND (status#5 = Completed))
    +- LogicalRDD [order_id#0, city#1, product#2, price#3, order_date#4, st

== Physical Plan ==
AdaptiveSparkPlan isFinalPlan=true
+- == Final Plan ==
  ResultQueryStage 1
    +- *(2) HashAggregate(keys=[city#26, order_id#0, price#47, product#27, ord
      +- AQEShuffleRead coalesced
        +- ShuffleQueryStage 0
          +- Exchange hashpartitioning(city#26, order_id#0, price#47, produ
            +- *(1) HashAggregate(keys=[city#26, order_id#0, price#47, pro
              +- *(1) Project [order_id#0, lower(trim(city#1, None)) AS c
                +- *(1) Filter (isnotnull(status#5) AND (status#5 = Comp
                  +- *(1) Scan ExistingRDD[order_id#0,city#1,product#2,
+- == Initial Plan ==
  HashAggregate(keys=[city#26, order_id#0, price#47, product#27, order_date#
    +- Exchange hashpartitioning(city#26, order_id#0, price#47, product#27, or
      +- HashAggregate(keys=[city#26, order_id#0, price#47, product#27, order_
        +- Project [order_id#0, lower(trim(city#1, None)) AS city#26, lower(
          +- Filter (isnotnull(status#5) AND (status#5 = Completed))
            +- Scan ExistingRDD[order_id#0,city#1,product#2,price#3,order_

```

Start coding or generate with AI.

